

Lecture 13: Decision Trees and Random Forests

Machine Learning, Summer Term 2019


July 1, 2019

Michael Tangermann Frank Hutter Marius Lindauer

University of Freiburg



The Big Picture



- Lecture 1: overview
- Lecture 2-6: linear methods
- Lecture 7-9: algorithm-independent principles
- Lectures 10-15: nonlinear methods
 - Lecture 10-12: kernel-based methods
 - Lectures 13-14: **tree-based methods and ensembles**
 - Lecture 15: neural networks

- 1 Decision and Regression Trees
 - Regression Trees
 - Classification Trees (= Decision Trees)
- 2 Bagging
- 3 Random Forests

1 Decision and Regression Trees

- Regression Trees
- Classification Trees (= Decision Trees)

2 Bagging

3 Random Forests

Motivation for Trees and Forests

For many applications, random forests are the **best off-the-shelf model**

- Trees: easy to **interpret**
- Directly handle **categorical features**
- Scalable to **many data points** (they are fast)
- Scalable to **many features** (they are automated feature selectors)
- Random forests: **robust** performance even for **small datasets**
- Random forests: **robust** to their **hyperparameter settings**
 - In contrast to, e.g., SVMs or neural networks


Trees and Forests are Extremely Popular Models



Leo Breiman 1928-2005

Classification and Regression Trees


Authors Leo Breiman, Jerome H Friedman, Richard A Olshen, Charles J Stone
Publication date 1999/5
Publisher CRC Press, New York
Total citations [Cited by 34504](#)



Scholar articles [Random forests](#)
L Breiman - Machine learning, 2001
[Cited by 29053](#) - [Related articles](#) - [All 68 versions](#)

Acknowledgement

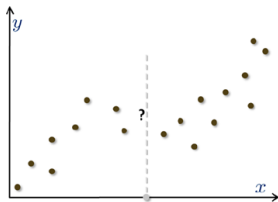
Most visualizations in this lecture are taken from this excellent book by Criminisi et. al (2013):




PDF of entire book available from university machines:
<http://link.springer.com/book/10.1007/978-1-4471-4929-3>

- 1 Decision and Regression Trees
 - Regression Trees
 - Classification Trees (= Decision Trees)
- 2 Bagging
- 3 Random Forests


Regression Trees




Idea: fit simple model to subset of the data



probab. model





polynomial



constant (standard)

We will only cover the standard case of constant leaf predictions

How to Split the Data

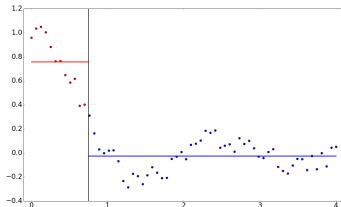


- Constant models in the leafs: \hat{f}_{left} and \hat{f}_{right}
- Greedily minimize sum of squared errors in the two children

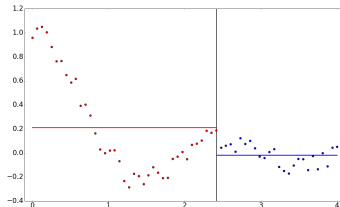
$$x_{\text{split}} \in \arg \min_x \left(\sum_{x_i \leq x} (y_i - \hat{f}_{left})^2 + \sum_{x_i > x} (y_i - \hat{f}_{right})^2 \right)$$

- Recursively split subsets

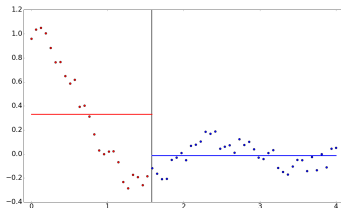
Let's vote: Which of These Splits is the Best?



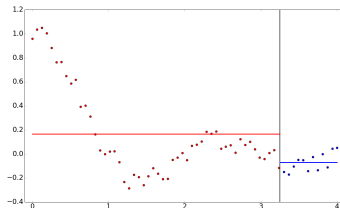
★ (black)



★ (red)

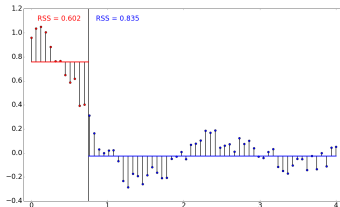


★ (yellow)

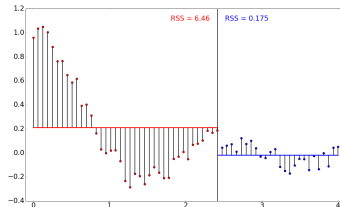


★ (blue)

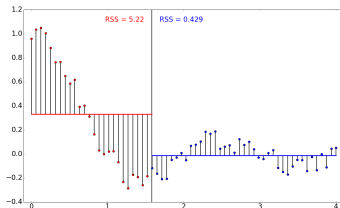
Let's vote: Which of These Splits is the Best?



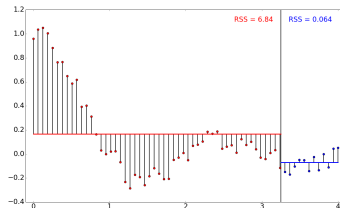
★ (black)



★ (red)



★ (yellow)



★ (blue)

CART algorithm [Breiman et. al 1984]

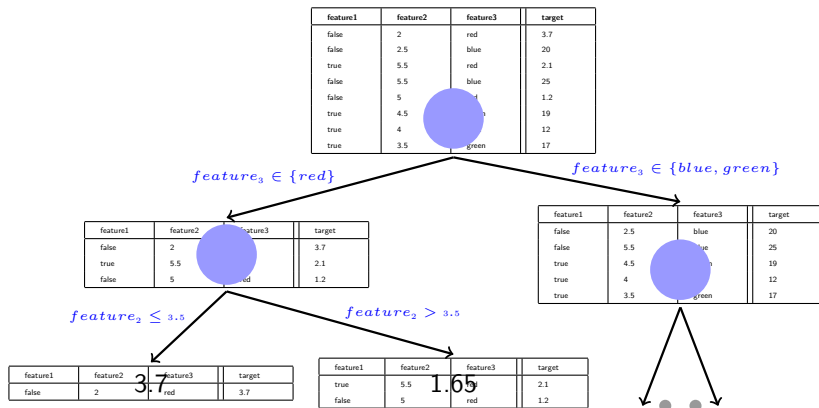
CART = Classification And Regression Trees

- Input: \mathbf{X} , \mathbf{y} (and hyperparameters `max_depth`, `min_leaf`)
- Check whether data should be split further; otherwise return leaf node
- Find best split value for each feature
- Choose best combination of split feature and split value
- Split data into left and right accordingly: $(\mathbf{X}_l, \mathbf{y}_l)$ and $(\mathbf{X}_r, \mathbf{y}_r)$
- Save split feature and value, and pointer to two new subtrees to be built recursively:

(CART($\mathbf{X}_l, \mathbf{y}_l, \text{max_depth}-1, \text{min_leaf}$),
CART($\mathbf{X}_r, \mathbf{y}_r, \text{max_depth}-1, \text{min_leaf}$))

Visualization of CART Algorithm For Regression: Training

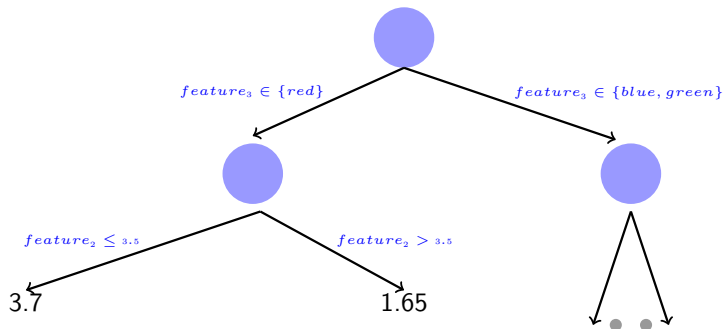
- In each internal node: only store split criterion used
- In each leaf: store mean of targets





Visualization of CART Algorithm For Regression: Prediction for New Inputs

E.g $x_{n+1} = (true, 4.7, red)$

- Walk down tree, return mean target stored in leaf $\Rightarrow 1.65$




Hierarchical Binary Splits



- Hierarchical splits (as on the right side) are easy to represent
 - Splits as in the left figure would be harder
- We could also use k -ary splits
 - But every k -ary split can be seen as a sequence of binary splits
 - Binary splits are faster and often yield better predictions


Formal Notation for Tree Predictions (1/2)

After the intuitive treatment so far, we now write down tree predictions formally.



- Trees partition the input space \mathcal{X} into regions R_1, \dots, R_J associated with their leaves
- Each leaf j has a simple model; here the constant γ_j
- Mathematically, a decision/regression tree T with constant leaf predictions is fully specified by $\langle R_1, \dots, R_J, \gamma_1, \dots, \gamma_J \rangle$

Formal Notation for Tree Predictions (2/2)



Tree Prediction

The prediction of a decision/regression tree with parameters $\Theta = \langle R_1, \dots, R_J, \gamma_1, \dots, \gamma_J \rangle$ is


$$T(x_i, \Theta) = \sum_{j=1}^J \gamma_j \mathbb{I}(x \in R_j)$$

Here and throughout, \mathbb{I} is the **indicator function**

$$\mathbb{I}(a) := \begin{cases} 1 & , \text{ if } a = \text{true} \\ 0 & , \text{ otherwise} \end{cases}$$

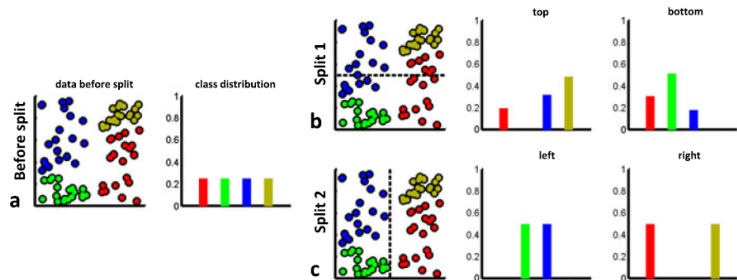
- 1 Decision and Regression Trees
 - Regression Trees
 - Classification Trees (= Decision Trees)
- 2 Bagging
- 3 Random Forests

Classification Trees (= Decision Trees)



- Conceptually the same as for regression
- Leaf model: majority vote; **probability of data in the leaf**
- Split criterion: Gini index; variance reduction; **information gain**

Classification Splitting Criteria




Intuitively, which of these splits is better?

- ★ Split 1 (into top and bottom)
- ★ Split 2 (into left and right)

Definition: Entropy

The **entropy** of a discrete random variable V with K possible outcomes v_k of respective probability $p(v_k)$ (for $k = 1, \dots, K$) is

$$H(V) = - \sum_{k=1}^K p(v_k) \log_2 p(v_k)$$



Examples for a Boolean random variable V with outcomes v_1 and v_2 :

- $p(v_1) = 0.5, p(v_2) = 0.5$. Then,
 $H(V) = -0.5 \log_2(0.5) - 0.5 \log_2(0.5) = 1$
- $p(v_1) = 0, p(v_2) = 1$. Then,
 $H(V) = -\lim_{x \rightarrow 0} x \log_2(x) - 1 \log_2(1) = 0$

Prominent Splitting Criterion: Information Gain

Definition: Information Gain

Let V denote a random variable with the empirical class distribution p of the N data points at the current node. Let a split s result in two child nodes with N_l and N_r data points and empirical class distributions p_l and p_r , and let V_l and V_r denote random variables with these probability distributions, respectively. Then, the **information gain** achieved by split s is:

$$\begin{aligned} I &= N \cdot H(V) - N_l \cdot H(V_l) - N_r \cdot H(V_r) \\ &= -N \sum_{k=1}^K p(v_k) \log_2(p(v_k)) + N_l \sum_{k=1}^K p_l(v_k) \log_2(p_l(v_k)) + N_r \sum_{k=1}^K p_r(v_k) \log_2(p_r(v_k)) \end{aligned}$$

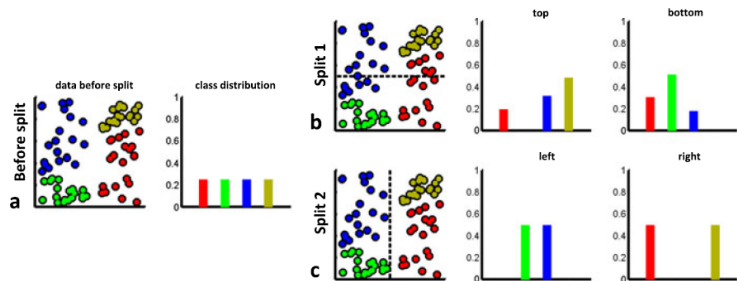
We want to maximize this information gain.

Information Gain: Example

$$\begin{aligned} I &= N \cdot H(V) - N_l \cdot H(V_l) - N_r \cdot H(V_r) \\ &= -N \sum_{k=1}^K p(v_k) \log_2(p(v_k)) + N_l \sum_{k=1}^K p_l(v_k) \log_2(p_l(v_k)) + N_r \sum_{k=1}^K p_r(v_k) \log_2(p_r(v_k)) \end{aligned}$$


- Example: Splitting 8 data points ($4 \times \text{false}$, $4 \times \text{true}$).
 - Split 1: left: $4 \times \text{true}$; right: $4 \times \text{false}$
 - $N \cdot H(V) = 8 \cdot 1 = 8$
 - $N_l \cdot H(V_l) = 4 \cdot 0 = 0$
 - $N_r \cdot H(V_r) = 4 \cdot 0 = 0$
 - Split 2: left: $2 \times \text{true}$, $2 \times \text{false}$; right: $2 \times \text{true}$, $2 \times \text{false}$
 - $N \cdot H(V) = 8 \cdot 1 = 8$
 - $N_l \cdot H(V_l) = 4 \cdot 1 = 4$
 - $N_r \cdot H(V_r) = 4 \cdot 1 = 4$

Information Gain: Example






- We previously said split 2 (into left and right) is intuitively better than split 1 (into top and bottom)
- Information gain quantifies this

Splits Along Several Dimensions



Axis-aligned splits are standard, but more complex splits are possible



Pros and Cons of Decision and Regression Trees

- + Flexible framework with exchangeable components:
splitting criterion, leaf model, type of split
- + Interpretability
- + Handle categorical input values natively
- + Handle unimportant features well
- + Scalable for large datasets
 - Tend to overfit
 - Deterministic, i.e. not suitable for some ensemble methods

Hyperparameters of Decision and Regression Trees

Regression and decision trees have several hyperparameters:

- Minimum number of samples in a leaf (`min_leaf`)
- Maximal depth of the tree (`max_depth`)
- Total number of nodes
- Leaf model (weak learner; here constant)
- Split criterion

Assignment 8 explores some aspects of their influence on the predictive quality.

Bias and Variance of Trees

- Facts about tree-based models
 - Trees are very expressive models
 - When you slightly change the data, you might get a very different tree
- Using these facts, please choose the right answer:
 - ★ Trees are a high-variance model.
 - ★ Trees are a low-variance model.
- Using these facts, please choose the right answer:
 - ★ Trees are a high-bias model.
 - ★ Trees are a low-bias model.

Reminder: Bias and Variance

$$\mathbb{E}_{\mathcal{D}}[\{\hat{h}(\mathbf{x}; \mathcal{D}) - h^*(\mathbf{x})\}^2] = \underbrace{\{\bar{h}(\mathbf{x}) - h^*(\mathbf{x})\}^2}_{(\text{bias})^2} + \underbrace{\mathbb{E}_{\mathcal{D}}[\{\hat{h}(\mathbf{x}; \mathcal{D}) - \bar{h}(\mathbf{x})\}^2]}_{\text{variance}}$$

expected squared deviation from best prediction = (bias)² + variance

- $\hat{h}(\mathbf{x}; \mathcal{D})$ is the prediction of the model that fits data \mathcal{D} best
- $h^*(\mathbf{x})$ is the true best (unknown) prediction
- $\bar{h}(\mathbf{x}) = \mathbb{E}_{\mathcal{D}}[\hat{h}(\mathbf{x}; \mathcal{D})]$ is the average model prediction
(average of models trained on data sets \mathcal{D} from a data distribution)

Computational Complexity of Regression and Decision Trees

- N data points of dimensionality D , axis-aligned splits
- Finding the **best split value for a given feature (pre-sorted)**: $O(N)$
 - Amortized analysis: $O(N)$ to initialize all data points to the left child, $O(1)$ for moving one data point from left to right at a time
- Finding best split point at root: $O(DN)$
- Let $T(N)$ denote the time we pay for a complete subtree with N data points; this has a part due to the cost at the root and parts due to the 2 smaller child subtrees
- **Best case**: balanced trees
 - Fitting: $T(N) = O(DN) + 2T(N/2)$
 - ↪ This leads to $O(DN \log N)$ since we pay $O(D \cdot N)$ at each of $\log N$ levels
 - Prediction: $O(\log N)$

Computational Complexity of Regression and Decision Trees

- N data points of dimensionality D , axis-aligned splits
- Finding the **best split value for a given feature (pre-sorted)**: $O(N)$
 - Amortized analysis: $O(N)$ to initialize all data points left to the left child, $O(1)$ for moving one data point from left to right at a time
- Finding best split point at root: $O(DN)$
- Let $T(N)$ denote the time we pay for a complete subtree with N data points; this has a part due to the cost at the root and parts due to the 2 smaller child subtrees
- **Worst case**: splitting off one data point at a time
 - Fitting: $T(N) = O(DN) + T(N - 1)$; this leads to $O(DN^2)$
 - Prediction: $O(N)$

- 1 Decision and Regression Trees
 - Regression Trees
 - Classification Trees (= Decision Trees)
- 2 Bagging
- 3 Random Forests


Bagging is a Committee / Ensemble Approach

- A single expert may fail
- A committee of experts is more likely to get it right (if experts are experienced and diverse)




Bagging [Breimann, 1996]

- Train N models on **bootstrap samples** of the training data
- For each model, data is drawn randomly with replacements
- **Average** output of all models (**bagging = bootstrap aggregation**)




Boostrapping


original data







histogram



resampled data




associated histograms



⋮

⋮

histogram of the means: bootstrap distribution of the mean



Why Does Bagging Work?

- Bagging
 - Train N models on **bootstrap samples** of the training data
 - For each model, data is drawn randomly with replacements
 - **Average** output of all models (**bagging = bootstrap aggregation**)

Discuss with your neighbor first.   (2 minutes)

- Voting question 1: A bagged estimator has
 - ★ higher variance
 - ★ lower variance

than the individual models it bags.

- Voting question 2: A bagged estimator has
 - ★ higher bias
 - ★ lower bias

than the individual models it bags.

Reminder: Bias and Variance

$$\mathbb{E}_{\mathcal{D}}[\{\hat{h}(\mathbf{x}; \mathcal{D}) - h^*(\mathbf{x})\}^2] = \underbrace{\{\bar{h}(\mathbf{x}) - h^*(\mathbf{x})\}^2}_{(\text{bias})^2} + \underbrace{\mathbb{E}_{\mathcal{D}}[\{\hat{h}(\mathbf{x}; \mathcal{D}) - \bar{h}(\mathbf{x})\}^2]}_{\text{variance}}$$


expected squared deviation from best prediction = (bias)² + variance

- $\hat{h}(\mathbf{x}; \mathcal{D})$ is the prediction of the model that fits data \mathcal{D} best
- $h^*(\mathbf{x})$ is the true best (unknown) prediction
- $\bar{h}(\mathbf{x}) = \mathbb{E}_{\mathcal{D}}[\hat{h}(\mathbf{x}; \mathcal{D})]$ is the average model prediction
(average of models trained on data sets \mathcal{D} from a data distribution)


- 1 Decision and Regression Trees
 - Regression Trees
 - Classification Trees (= Decision Trees)
- 2 Bagging
- 3 Random Forests

Why Does Bagging Work Well for Trees?

- Theoretical result [Breiman 2001]:
 - Ensemble error depends on (1) **strength** of individual models and (2) the degree to which the models' errors are **uncorrelated**.
- We thus want:
 - Individual models that still work (quite) well
 - Quite different models
- **Randomized** trees give us these properties. We can randomize in many ways:
 - best split using a random subset of $m \leq D$ features
 - splitting using best out of a fixed number of random splits
 - training on bootstrap-samples of the data

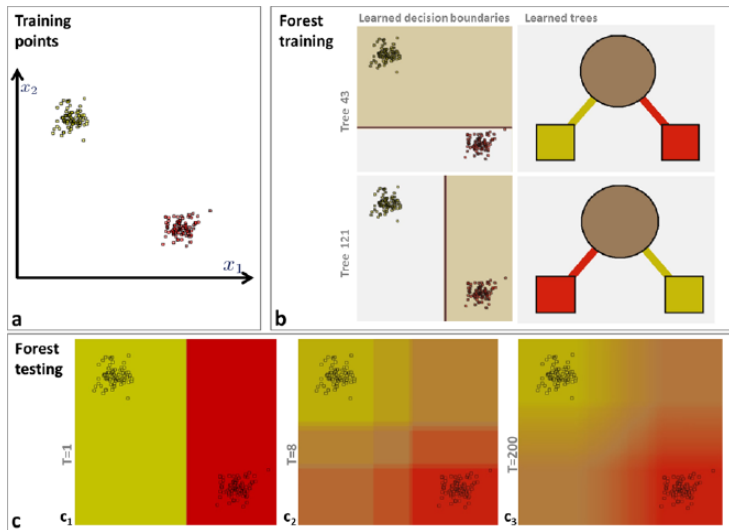


a) Low randomness, high tree correlation



b) High randomness, low tree correlation

Visualization of Random Forests and Their Predictions



Algorithm to Grow a Random Forest

Algorithm 15.1 *Random Forest for Regression or Classification.*

1. For $b = 1$ to B :
 - (a) Draw a bootstrap sample \mathbf{Z}^* of size N from the training data.
 - (b) Grow a random-forest tree T_b to the bootstrapped data, by recursively repeating the following steps for each terminal node of the tree, until the minimum node size n_{min} is reached.
 - i. Select m variables at random from the p variables.
 - ii. Pick the best variable/split-point among the m .
 - iii. Split the node into two daughter nodes.
2. Output the ensemble of trees $\{T_b\}_1^B$.

To make a prediction at a new point x :

Regression: $\hat{f}_{\text{rf}}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$.

Classification: Let $\hat{C}_b(x)$ be the class prediction of the b th random-forest tree. Then $\hat{C}_{\text{rf}}^B(x) = \text{majority vote } \{\hat{C}_b(x)\}_1^B$.

from [Hastie, Tibshirani and Friedman]

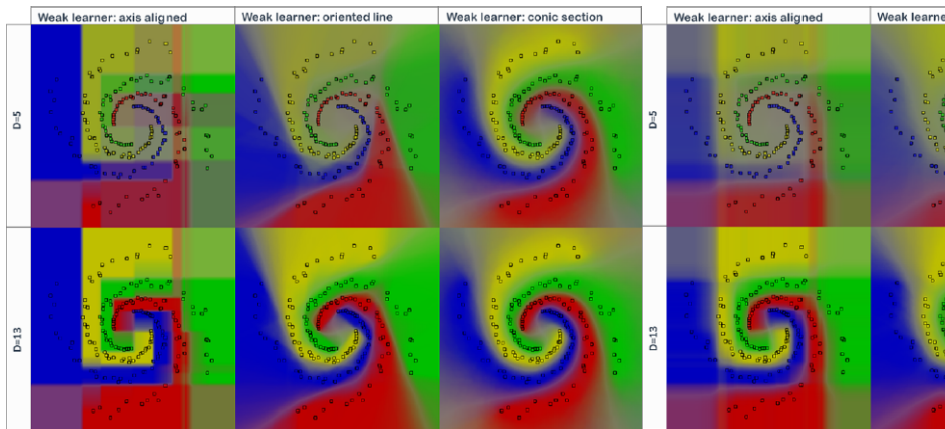
Computational Complexity of Random Forests

- N data points of dimensionality D , axis-aligned splits
- Picking m variables at random at each split; B trees
- **Best case:** balanced trees
 - Fitting: $O(BmN \log N)$
 - Prediction: $O(B \log N)$
- **Worst case:** splitting off one data point at a time
 - Fitting: $O(BmN^2)$
 - Prediction: $O(BN)$

Some Ways of Ensembling Trees

- Decision Trees + Bagging = Bagged trees
- Decision Trees + random feature subsets + Bagging = Random forest
- Decision Trees with almost random splits + Bagging = ExtraTrees
 - Extremely randomized trees

Examples: Influence of split randomization



- 400 trees with `max_depth` 5 (top) and 13 (bottom)
- Different split types: axis aligned (left), oriented lines (middle), and conic sections (right)
- Splitting: best out of 500 random proposed splits for each splits

Advantages of Boostapping

- Can help detect outliers
- Decorrelation of the trees in the ensemble
- Out-of-bag error:
 - not every data point is used to fit every single tree
 - in fact, almost 37% are not used in each tree (see assignment)
 - predict unused points for each tree to obtain unbiased estimate of the generalization error

Pros and Cons of Random Forests (and co.)

- + All pros from decision trees remain (except interpretability)
- + Better generalization
- + Out-of-bag error with little overhead
- + Scalability to large data sets and high dimensions
- + Require little tuning
 - Relatively weak performance for smooth functions without noise

Summary by learning goals

Having heard this lecture, you can now . . .

- determine **good splits** in regression and classification trees
- describe the steps of the **CART algorithm**
- describe the steps of the **random forest algorithm**
- formally describe **entropy** and **information gain**
- derive the **complexity** of decision trees and random forests
- explain why **bootstrapping** works well **for trees**
- describe some ways to **randomize trees** and their effects

Further Reading

- Criminisi et. al (2013)
 - Chapter 3: Introduction: The Abstract Forest Model
 - Chapter 4: Classification Forests
 - Chapter 5: Regression Forests
 - PDF of entire book available from university machines: <http://link.springer.com/book/10.1007/978-1-4471-4929-3>
- Hastie, Tibshirani and Friedman
 - Section 9.2: Tree-Based Methods
 - Chapter 15: Random Forests



Preview of Assignment 8

In assignment 8, you will ...

- implement a simple regression tree and forest
- study hyperparameter influence on toy data
- calculate entropy and information gain by hand
- build a small decision tree by hand